

ENGG 4030 Homework_2

I declare that the assignment submitted on Elearning system is original except for source material explicitly acknowledged, and that the same or related material has not been previously submitted for another course. I also acknowledge that I am aware of University policy and regulations on honesty in academic work, and of the disciplinary guidelines and procedures applicable to breaches of such policy and regulations, as contained in the website <http://www.cuhk.edu.hk/policy/academichonesty/> .

Signed (Student 柳一村) Date: 18/03/2018
Name LIU Yicun SID 1155092202

Q1

Pre-processing:

Due to the fact that we have two datasets “shakespeare1” and “shakespeare2” in hand, I simply combine the two input files as “shakespeare3” and use it as input of the following task.

Note that in the second original dataset “shakespeare2”, there are some duplicate words in each line, here I choose to delete them in the mapper, using data structure named *Set* in python:

```
words = line.split()
words = list(set(words))
```

A: Find Frequent Item Sets

- (a) As the requirement said, each line means a basket, and here we need to implement a non-distribute version of A-priori to find the top-40 frequent pairs.

In the first pass, we count the frequency of single item and find the frequent items:

```
with open(fname) as f:
    lines = [line.rstrip('\n') for line in open(fname)]
    for line in lines:
        num_lines += 1
        words = line.split()
        words = list(set(words))
        for word in words:
            if word not in wordfreq:
                wordfreq[word] = 1
            else:
                wordfreq[word] += 1

    for word in wordfreq:
        if (wordfreq[word] >= threshold * num_lines):
            temp[word] = wordfreq[word]
```

In the second pass, I read baskets again and count in main memory only the pairs where both elements are frequent, and output the final pairs with top-40 frequencies:

```
# Second step of A-Priori
with open(fname) as f:
    lines = [line.rstrip('\n') for line in open(fname)]
    for line in lines:
        words = line.split()
        words = list(set(words))
        for j in range(0, len(words)-1):
            for k in range(j+1, len(words)):
                if (words[j] in wordfreq) and (words[k] in wordfreq):
                    if (words[j] <= words[k]):
                        pair = words[j] + "," + words[k]
                    else:
                        pair = words[k] + "," + words[j]
                    if pair not in pairfreq:
                        pairfreq[pair] = 1
                    else:
                        pairfreq[pair] += 1

    if (len(pairfreq) > top_num):
        topkeys = sorted(pairfreq, key=pairfreq.get, reverse=True)[:top_num]
    else:
        topkeys = sorted(pairfreq, key=pairfreq.get, reverse=False)

    for topkey in topkeys:
        print(topkey, pairfreq[topkey])
```

Consider it is a non-distributed version, I simply run the code on my MacBook (late 2015, with core-m3), the result is:

('of,the', 352310)	('a,to', 57495)	('be,to', 40461)
('and,the', 174876)	('a,in', 54897)	('as,the', 39145)
('in,the', 146017)	('is,the', 54866)	('of,was', 36443)
('the,to', 145171)	('on,the', 53910)	('from,the', 34846)
('and,of', 106824)	('that,the', 51819)	('and,was', 34740)
('a,of', 100121)	('thee,thou', 50082)	('and,his', 34170)
('a,the', 84377)	('by,the', 47365)	('his,of', 33608)
('and,to', 74266)	('the,with', 44846)	('it,to', 33164)
('in,of', 73828)	('he,the', 44370)	('of,that', 33117)
('of,to', 68015)	('for,the', 43644)	('his,the', 33109)
('a,and', 66899)	('it,the', 41316)	('art,thou', 32978)
('and,in', 62592)	('thee,thy', 41252)	('i,the', 32738)
('the,was', 62453)	('at,the', 40688)	
('thou,thy', 61455)	('in,to', 40610)	

The running time is 4 mins 10 secs.

- (b) Here I need to implement a distributed version of SON. Note that the first pass of SON is mainly A-Priori, so that I can utilize the code from part (a):

```
p = 1
num_lines = 0
threshold = 0.005

for line in sys.stdin:
    num_lines += 1
    line = line.strip()

    basket.append(line)

    words = line.split()
    words = list(set(words))
    for word in words:
        if word not in wordfreq:
            wordfreq[word] = 1
        else:
            wordfreq[word] += 1

for word in wordfreq:
    if (wordfreq[word] >= threshold * num_lines / p):
        temp[word] = wordfreq[word]

wordfreq = copy.deepcopy(temp)
temp = {}

for i in range(0, len(basket)):
    line = basket[i]
    word_basket = line.strip().split()
    word_basket = list(set(word_basket))
    for j in range(0, len(word_basket)-1):
        for k in range(j+1, len(word_basket)):
            if (word_basket[j] in wordfreq):
                if (word_basket[k] in wordfreq):
                    if (word_basket[j] <= word_basket[k]):
                        pair = word_basket[j] + "," + word_basket[k]
                    else:
                        pair = word_basket[k] + "," + word_basket[j]
                    if pair not in pairfreq:
                        pairfreq[pair] = 1
                    else:
                        pairfreq[pair] += 1

for pair in pairfreq:
    if (pairfreq[pair] >= threshold * num_lines / p):
        print '%s\t%s' % (pair, str(pairfreq[pair]))
```

```
import sys

current_pair = None
for line in sys.stdin:
    pair, count = line.strip().split('\t')
    if pair == current_pair:
        continue
    else:
        if current_pair:
            print current_pair
            current_pair = pair
        if current_pair == pair:
            print current_pair
```

Mapper / Reducer in the first MapReduce job

In the first MapReduce job, I use A-Priori find the candidate pairs, which means that they should be frequent in at least one input file. The reducer in the first job stores the candidate pair into an intermediate file named "candidatepair_b.txt".

```
import sys

fname = 'candidatepair_b.txt'

threshold = 0.005

pairfreq = {}

with open(fname) as f:
    lines = [line.rstrip('\n') for line in open(fname)]
    for line in lines:
        pair = line.strip()
        pairfreq[pair] = 0

for line in sys.stdin:
    line = line.strip()
    words = line.split()
    words = list(set(words))
    for j in range(0, len(words)-1):
        for k in range(j+1, len(words)):
            if words[j] <= words[k]:
                pair = words[j] + "," + words[k]
            else:
                pair = words[k] + "," + words[j]
            if pair in pairfreq:
                pairfreq[pair] += 1

for pair in pairfreq:
    print '%s\t%s' % (pair, str(pairfreq[pair]))
```

```
import sys

pairfreq = {}

threshold = 0.005
p = 1
num_lines = 4340061

current_pair = None
for line in sys.stdin:
    pair, count = line.strip().split("\t")
    count = int(count)
    if current_pair == pair:
        total += count
    else:
        if current_pair:
            if total >= num_lines * threshold / p:
                print current_pair + "\t" + str(total)
            current_pair = pair
        total = count
```

Mapper / Reducer in the second MapReduce job

In the second MapReduce job, the mapper should read the candidate pairs from the intermediate file and counts only the candidate frequent pairs, while the reducer should combine and output the real frequent pairs. Note that when the reducer number is more than zero, the total threshold to determine the true frequent pairs should be divided by p (the number of reducers). I integrate all the commands into a bash file:

```
chmod +x mapper_b1.py mapper_b2.py reducer_b1.py reducer_b2.py

hdfs dfs -rm -r /user/ly116/output_b1
hdfs dfs -rm -r /user/ly116/output_b2

hdfs dfs -mkdir /user/ly116/shakespeare3
hdfs dfs -put ./shakespeare_basket/shakespeare_basket3 /user/ly116/shakespeare3

hadoop jar /usr/hdp/current/hadoop-mapreduce-client/hadoop-streaming.jar \
-D mapreduce.map.memory.mb=4096 \
-D mapreduce.reduce.memory.mb=4096 \
-D mapred.map.tasks=20 \
-D mapred.reduce.tasks=1 \
-input /user/ly116/shakespeare3/* \
-output output_b1 \
-file mapper_b1.py -mapper mapper_b1.py \
-file reducer_b1.py -reducer reducer_b1.py

hdfs dfs -getmerge output_b1 ~/candidatepair_b.txt

hadoop jar /usr/hdp/current/hadoop-mapreduce-client/hadoop-streaming.jar \
-D mapreduce.map.memory.mb=4096 \
-D mapreduce.reduce.memory.mb=4096 \
-D mapred.map.tasks=20 \
-D mapred.reduce.tasks=1 \
-input /user/ly116/shakespeare3/* \
-output output_b2 \
-file mapper_b2.py -mapper mapper_b2.py \
-file reducer_b2.py -reducer reducer_b2.py \
-file candidatepair_b.txt

hdfs dfs -cat /user/ly116/output_b2/* > out_b.txt

hdfs dfs -rm -r /user/ly116/output_b1
hdfs dfs -rm -r /user/ly116/output_b2
```

First MapReduce: 20 mappers, 1 reducer, avg execution time is 14s.

Job Name:	streamjob587191504991495096.jar
User Name:	ly116
Queue:	default
State:	SUCCEEDED
Uberized:	false
Submitted:	Mon Mar 19 01:45:00 HKT 2018
Started:	Mon Mar 19 01:45:05 HKT 2018
Finished:	Mon Mar 19 01:45:20 HKT 2018
Elapsed:	14sec
Diagnostics:	
Average Map Time	9sec
Average Shuffle Time	3sec
Average Merge Time	0sec
Average Reduce Time	0sec

ApplicationMaster			
Attempt Number	Start Time	Node	Logs
1	Mon Mar 19 01:45:03 HKT 2018	dic13.ie.cuhk.edu.hk:8042	logs

Task Type	Total	Complete
Map	20	20
Reduce	1	1

Second MapReduce: 20 mappers, 1 reducer, avg execution time is 15s.

		Job Overview
Job Name:	streamjob8264891015266369938.jar	
User Name:	ly116	
Queue:	default	
State:	SUCCEEDED	
Uberized:	false	
Submitted:	Mon Mar 19 01:45:26 HKT 2018	
Started:	Mon Mar 19 01:45:31 HKT 2018	
Finished:	Mon Mar 19 01:45:46 HKT 2018	
Elapsed:	15sec	
Diagnostics:		
Average Map Time	9sec	
Average Shuffle Time	3sec	
Average Merge Time	0sec	
Average Reduce Time	0sec	

ApplicationMaster			
Attempt Number	Start Time	Node	Logs
1	Mon Mar 19 01:45:29 HKT 2018	dic14.ie.cuhk.edu.hk:8042	logs

Task Type	Total	Complete
Map	20	20
Reduce	1	1

Benefited from distributed structure of multiple mappers and reducers, the overall execution time is only about 1/9 of the execution time of (a).

The final result (after sorting the top 40) is:

```

1 of,the 352310
2 and,the 174876
3 in,the 146017
4 the,to 145171
5 and,of 106824
6 a,of 100121
7 a,the 84377
8 and,to 74266
9 in,of 73828
10 of,to 68015
11 a,and 66899
12 and,in 62592
13 the,was 62453
14 thou,thy 61455
15 a,to 57495
16 a,in 54897
17 is,the 54866
18 on,the 53910
19 that,the 51819
20 thee,thou 50082
21 by,the 47365
22 the,with 44846
23 he,the 44370
24 for,the 43644
25 it,the 41316
26 thee,thy 41252
27 at,the 40688
28 in,to 40610
29 be,to 40461
30 as,the 39145
31 of,was 36443
32 from,the 34846
33 and,was 34740
34 and,his 34170
35 his,of 33608
36 it,to 33164
37 of,that 33117
38 his,the 33109
39 art,thou 32978
40 i,the 32738

```

- (c) For the SON of tuples, the main difference is the first mapper, which is expected to output the candidate tuples, here we add 'find candidate tuples' to the previous A-Priori code for pairs. With frequent pairs in hand, I read baskets for the third time and count in main memory only the tuples where six possible pairs of it are all frequent:

```

for i in range(0, len(basket)):
    line = basket[i]
    word_basket = line.strip().split()
    word_basket = list(set(word_basket))
    for j in range(0, len(word_basket) - 2):
        for k in range(j+1, len(word_basket) - 1):
            for l in range(k+1, len(word_basket)):
                if(word_basket[j] <= word_basket[k]):
                    pair_1 = word_basket[j] + "," + word_basket[k]
                else:
                    pair_1 = word_basket[k] + "," + word_basket[j]
                if(word_basket[j] <= word_basket[l]):
                    pair_2 = word_basket[j] + "," + word_basket[l]
                else:
                    pair_2 = word_basket[l] + "," + word_basket[j]
                if(word_basket[k] <= word_basket[l]):
                    pair_3 = word_basket[k] + "," + word_basket[l]
                else:
                    pair_3 = word_basket[l] + "," + word_basket[k]

                if(pair_1 in pairfreq) and (pair_2 in pairfreq) and (pair_3 in pairfreq):
                    if(word_basket[j] < word_basket[k]) and (word_basket[k] < word_basket[l]):
                        tupl = word_basket[j] + "," + word_basket[k] + "," + word_basket[l]
                    if(word_basket[j] < word_basket[l]) and (word_basket[l] < word_basket[k]):
                        tupl = word_basket[j] + "," + word_basket[l] + "," + word_basket[k]
                    if(word_basket[k] < word_basket[l]) and (word_basket[l] < word_basket[j]):
                        tupl = word_basket[k] + "," + word_basket[l] + "," + word_basket[j]
                    if(word_basket[k] < word_basket[j]) and (word_basket[j] < word_basket[l]):
                        tupl = word_basket[k] + "," + word_basket[j] + "," + word_basket[l]
                    if(word_basket[l] < word_basket[j]) and (word_basket[j] < word_basket[k]):
                        tupl = word_basket[l] + "," + word_basket[j] + "," + word_basket[k]
                    if(word_basket[l] < word_basket[k]) and (word_basket[k] < word_basket[j]):
                        tupl = word_basket[l] + "," + word_basket[k] + "," + word_basket[j]
                    if tupl not in tuplefreq:
                        tuplefreq[tupl] = 1
                    else:
                        tuplefreq[tupl] += 1

```

First MapReduce: 40 mappers, 10 reducers, avg execution time is 1min 14s.

Job Overview			
Job Name: streamjob2253674052302093918.jar			
User Name: ly116			
Queue: default			
State: SUCCEEDED			
Uberized: false			
Submitted: Mon Mar 19 01:46:38 HKT 2018			
Started: Mon Mar 19 01:46:43 HKT 2018			
Finished: Mon Mar 19 01:47:57 HKT 2018			
Elapsed: 1mins, 14sec			
Diagnostics:			
Average Map Time 33sec			
Average Shuffle Time 57sec			
Average Merge Time 0sec			
Average Reduce Time 0sec			

ApplicationMaster			
Attempt Number	Start Time	Node	Logs
1	Mon Mar 19 01:46:41 HKT 2018	dic8.ie.cuhk.edu.hk:8042	logs

Task Type	Total	Complete
Map	40	40
Reduce	10	10

Second MapReduce: 40 mappers, 1 reducer, avg execution time is 1min 30s.

Job Overview			
Job Name:	streamjob7988286556542426190.jar		
User Name:	ly116		
Queue:	default		
State:	SUCCEEDED		
Uberized:	false		
Submitted:	Mon Mar 19 01:48:04 HKT 2018		
Started:	Mon Mar 19 01:48:08 HKT 2018		
Finished:	Mon Mar 19 01:49:38 HKT 2018		
Elapsed:	1mins, 30sec		
Diagnostics:			
Average Map Time	38sec		
Average Shuffle Time	1mins, 14sec		
Average Merge Time	0sec		
Average Reduce Time	0sec		

ApplicationMaster			
Attempt Number	Start Time	Node	Logs
1	Mon Mar 19 01:48:06 HKT 2018	dic12.ie.cuhk.edu.hk:8042	logs

Task Type	Total	Complete
Map	40	40
Reduce	1	1

The final result (after sorting the top 20) is:

```

1 and,of,the 57234
2 in,of,the 47475
3 of,the,to 40526
4 a,of,the 38061
5 and,in,the 25216
6 and,the,to 24450
7 act,enter,scene 23816
8 thee,thou,thy 21170
9 of,the,was 20549
10 act,exeunt,scene 20404
11 of,on,the 18828
12 is,of,the 17794
13 enter,exeunt,scene 17087
14 act,enter,exeunt 17080
15 a,in,the 16892
16 in,the,to 16381
17 a,and,of 16350
18 of,that,the 15906
19 a,and,the 15521
20 by,of,the 15396

```

- (d) To implement PCY on the basis of the previous SON construction, we need to maintain a hash table in the first mapper:

```
hashtable = [0 for x in range(100000)]
```

For each line of the first stdin, I hash every possible combination of pairs in that line into the hash table and add 1 to the count.

```

for j in range(0, len(words)-1):
    for k in range(j+1, len(words)):
        hash_val = hash(words[j] + words[k]) % 100000
        hashtable[hash_val] += 1

```


By the assumption that “only pairs in frequent buckets are frequent”, I check the count in the hash table in the second pass of A-Priori and only take the qualified ones as candidate pairs:

```
for i in range (0,len(basket)):
    line = basket[i]
    word_basket = line.strip().split()
    word_basket = list(set(word_basket))
    for j in range (0, len(word_basket)-1):
        for k in range(j+1,len(word_basket)):
            if(word_basket[j] in wordfreq):
                if(word_basket[k] in wordfreq):
                    hash_val = hash(word_basket[j] + word_basket[k]) % 100000
                    if(hashtable[hash_val] >= threshold * num_lines / p):
                        if(word_basket[j]<= word_basket[k]):
                            pair = word_basket[j]+","+ word_basket[k]
                        else:
                            pair = word_basket[k]+","+ word_basket[j]
                        if pair not in pairfreq:
                            pairfreq[pair] = 1
                        else:
                            pairfreq[pair] += 1
```

The output is the same as (b), which ensure our implementation output the correct answer. The running time for (b), (c), (d) is listed below:

Algorithm	1st MapReduce	2th MapReduce	Total	Intermediate File Size
(a)	NA	NA	4min 10s	NA
(b)	14s	15s	29s	96.1 KB
(c)	1 min 14s	1 min 30s	2min 44s	1.4 MB
(d)	16s	12s	28s	70.3 KB

Comparison of (b) and (d): PCY takes longer time in the 1st MapReduce job to do the hashing but reduce the intermediate candidate pair file size by about 20%, which means that it only need to maintain a matrix/dictionary in the second MapReduce job and the execution time in the second job is shorter.

Comparison of (a) and (b): The distributed version (b) (20 mappers and 1 reducer) saves more than 80% time than the non-distributed version (a). Benefited from multiple mappers, it can output the candidate pairs in parallel and so is counting the real frequent pairs from the candidate pairs.

Comparison of (b) and (c): considering (b) need to find tuples but not pairs, which requires N times more search in each stdin line, another pass round of A-Priori in the first mapper and needs N times more RAM in the second MapReduce job, the execution time for both two jobs are significantly longer.

Q2 Locality-Sensitive Hashing

(a) $P1 = 1 - [1 - (T_1)^r]^B$ $P2 = 1 - [1 - (T_2)^r]^B$

(b) When $T1=0.9$, $T2=0.7$, $P1=0.99$ and $P2=0.05$, we can choose $r = 20$ while $B = 40$. to satisfy that circumstance.